

Logistic Regression

Logistic Regression: The Model

- A model for doing *probabilistic* binary classification
- Predicts *label probabilities* rather than a hard value of the label

$$\begin{aligned}p(y_n = 1 | \mathbf{x}_n, \mathbf{w}) &= \mu_n \\p(y_n = 0 | \mathbf{x}_n, \mathbf{w}) &= 1 - \mu_n\end{aligned}$$

Logistic Regression: The Model

- A model for doing *probabilistic* binary classification
- Predicts *label probabilities* rather than a hard value of the label

$$\begin{aligned}p(y_n = 1 | \mathbf{x}_n, \mathbf{w}) &= \mu_n \\p(y_n = 0 | \mathbf{x}_n, \mathbf{w}) &= 1 - \mu_n\end{aligned}$$

- The model's prediction is a probability defined using the **sigmoid function**

$$f(\mathbf{x}_n) = \mu_n = \sigma(\mathbf{w}^\top \mathbf{x}_n) = \frac{1}{1 + \exp(-\mathbf{w}^\top \mathbf{x}_n)} = \frac{\exp(\mathbf{w}^\top \mathbf{x}_n)}{1 + \exp(\mathbf{w}^\top \mathbf{x}_n)}$$

Logistic Regression: The Model


- A model for doing *probabilistic* binary classification
- Predicts *label probabilities* rather than a hard value of the label

$$\begin{aligned}p(y_n = 1 | \mathbf{x}_n, \mathbf{w}) &= \mu_n \\p(y_n = 0 | \mathbf{x}_n, \mathbf{w}) &= 1 - \mu_n\end{aligned}$$

- The model's prediction is a probability defined using the **sigmoid function**

$$f(\mathbf{x}_n) = \mu_n = \sigma(\mathbf{w}^\top \mathbf{x}_n) = \frac{1}{1 + \exp(-\mathbf{w}^\top \mathbf{x}_n)} = \frac{\exp(\mathbf{w}^\top \mathbf{x}_n)}{1 + \exp(\mathbf{w}^\top \mathbf{x}_n)}$$

- The sigmoid first computes a real-valued “score” $\mathbf{w}^\top \mathbf{x} = \sum_{d=1}^D w_d x_d$ and “squashes” it between (0,1) to turn this score into a **probability score**



Logistic Regression: The Model


- A model for doing *probabilistic* binary classification
- Predicts *label probabilities* rather than a hard value of the label

$$\begin{aligned}p(y_n = 1 | \mathbf{x}_n, \mathbf{w}) &= \mu_n \\p(y_n = 0 | \mathbf{x}_n, \mathbf{w}) &= 1 - \mu_n\end{aligned}$$

- The model's prediction is a probability defined using the **sigmoid function**

$$f(\mathbf{x}_n) = \mu_n = \sigma(\mathbf{w}^\top \mathbf{x}_n) = \frac{1}{1 + \exp(-\mathbf{w}^\top \mathbf{x}_n)} = \frac{\exp(\mathbf{w}^\top \mathbf{x}_n)}{1 + \exp(\mathbf{w}^\top \mathbf{x}_n)}$$

- The sigmoid first computes a real-valued “score” $\mathbf{w}^\top \mathbf{x} = \sum_{d=1}^D w_d x_d$ and “squashes” it between (0,1) to turn this score into a **probability score**



- Model parameter is the unknown \mathbf{w} . Need to learn it from training data.

Logistic Regression: An Interpretation

- Recall that the logistic regression model defines

$$\begin{aligned} p(y = 1 | \mathbf{x}, \mathbf{w}) &= \mu = \sigma(\mathbf{w}^\top \mathbf{x}) = \frac{1}{1 + \exp(-\mathbf{w}^\top \mathbf{x})} = \frac{\exp(\mathbf{w}^\top \mathbf{x})}{1 + \exp(\mathbf{w}^\top \mathbf{x})} \\ p(y = 0 | \mathbf{x}, \mathbf{w}) &= 1 - \mu = 1 - \sigma(\mathbf{w}^\top \mathbf{x}) = \frac{1}{1 + \exp(\mathbf{w}^\top \mathbf{x})} \end{aligned}$$

Logistic Regression: An Interpretation

- Recall that the logistic regression model defines

$$p(y = 1|\mathbf{x}, \mathbf{w}) = \mu = \sigma(\mathbf{w}^\top \mathbf{x}) = \frac{1}{1 + \exp(-\mathbf{w}^\top \mathbf{x})} = \frac{\exp(\mathbf{w}^\top \mathbf{x})}{1 + \exp(\mathbf{w}^\top \mathbf{x})}$$

$$p(y = 0|\mathbf{x}, \mathbf{w}) = 1 - \mu = 1 - \sigma(\mathbf{w}^\top \mathbf{x}) = \frac{1}{1 + \exp(\mathbf{w}^\top \mathbf{x})}$$

- The log-odds of this model

$$\log \frac{p(y = 1|\mathbf{x}, \mathbf{w})}{p(y = 0|\mathbf{x}, \mathbf{w})} = \log \exp(\mathbf{w}^\top \mathbf{x}) = \mathbf{w}^\top \mathbf{x}$$

Logistic Regression: An Interpretation

- Recall that the logistic regression model defines

$$\begin{aligned} p(y = 1|\mathbf{x}, \mathbf{w}) &= \mu = \sigma(\mathbf{w}^\top \mathbf{x}) = \frac{1}{1 + \exp(-\mathbf{w}^\top \mathbf{x})} = \frac{\exp(\mathbf{w}^\top \mathbf{x})}{1 + \exp(\mathbf{w}^\top \mathbf{x})} \\ p(y = 0|\mathbf{x}, \mathbf{w}) &= 1 - \mu = 1 - \sigma(\mathbf{w}^\top \mathbf{x}) = \frac{1}{1 + \exp(\mathbf{w}^\top \mathbf{x})} \end{aligned}$$

- The log-odds of this model

$$\log \frac{p(y = 1|\mathbf{x}, \mathbf{w})}{p(y = 0|\mathbf{x}, \mathbf{w})} = \log \exp(\mathbf{w}^\top \mathbf{x}) = \mathbf{w}^\top \mathbf{x}$$

- Thus if $\mathbf{w}^\top \mathbf{x} > 0$ then the positive class is more probable

Logistic Regression: An Interpretation


- Recall that the logistic regression model defines

$$\begin{aligned} p(y = 1|\mathbf{x}, \mathbf{w}) &= \mu = \sigma(\mathbf{w}^\top \mathbf{x}) = \frac{1}{1 + \exp(-\mathbf{w}^\top \mathbf{x})} = \frac{\exp(\mathbf{w}^\top \mathbf{x})}{1 + \exp(\mathbf{w}^\top \mathbf{x})} \\ p(y = 0|\mathbf{x}, \mathbf{w}) &= 1 - \mu = 1 - \sigma(\mathbf{w}^\top \mathbf{x}) = \frac{1}{1 + \exp(\mathbf{w}^\top \mathbf{x})} \end{aligned}$$

- The log-odds of this model

$$\log \frac{p(y = 1|\mathbf{x}, \mathbf{w})}{p(y = 0|\mathbf{x}, \mathbf{w})} = \log \exp(\mathbf{w}^\top \mathbf{x}) = \mathbf{w}^\top \mathbf{x}$$

- Thus if $\mathbf{w}^\top \mathbf{x} > 0$ then the positive class is more probable
- A linear classification model. Separates the two classes via a hyperplane (similar to other linear classification models such as Perceptron and SVM)



Loss Function Optimization View for Logistic Regression

Logistic Regression: The Loss Function

- What loss function to use? One option is to use the squared loss

$$\ell(y_n, f(\mathbf{x}_n)) = (y_n - f(\mathbf{x}_n))^2 = (y_n - \mu_n)^2 = (y_n - \sigma(\mathbf{w}^\top \mathbf{x}_n))^2$$

Logistic Regression: The Loss Function

- What loss function to use? One option is to use the squared loss

$$\ell(y_n, f(\mathbf{x}_n)) = (y_n - f(\mathbf{x}_n))^2 = (y_n - \mu_n)^2 = (y_n - \sigma(\mathbf{w}^\top \mathbf{x}_n))^2$$

- This is **non-convex** and not easy to optimize

Logistic Regression: The Loss Function

- What loss function to use? One option is to use the squared loss

$$\ell(y_n, f(\mathbf{x}_n)) = (y_n - f(\mathbf{x}_n))^2 = (y_n - \mu_n)^2 = (y_n - \sigma(\mathbf{w}^\top \mathbf{x}_n))^2$$

- This is **non-convex** and not easy to optimize
- Consider the following loss function

$$\ell(y_n, f(\mathbf{x}_n)) = \begin{cases} -\log(\mu_n) & y_n = 1 \\ -\log(1 - \mu_n) & y_n = 0 \end{cases}$$

Logistic Regression: The Loss Function

- What loss function to use? One option is to use the squared loss

$$\ell(y_n, f(\mathbf{x}_n)) = (y_n - f(\mathbf{x}_n))^2 = (y_n - \mu_n)^2 = (y_n - \sigma(\mathbf{w}^\top \mathbf{x}_n))^2$$

- This is **non-convex** and not easy to optimize
- Consider the following loss function

$$\ell(y_n, f(\mathbf{x}_n)) = \begin{cases} -\log(\mu_n) & y_n = 1 \\ -\log(1 - \mu_n) & y_n = 0 \end{cases}$$

- This loss function makes intuitive sense

Logistic Regression: The Loss Function

- What loss function to use? One option is to use the squared loss

$$\ell(y_n, f(\mathbf{x}_n)) = (y_n - f(\mathbf{x}_n))^2 = (y_n - \mu_n)^2 = (y_n - \sigma(\mathbf{w}^\top \mathbf{x}_n))^2$$

- This is **non-convex** and not easy to optimize
- Consider the following loss function

$$\ell(y_n, f(\mathbf{x}_n)) = \begin{cases} -\log(\mu_n) & y_n = 1 \\ -\log(1 - \mu_n) & y_n = 0 \end{cases}$$

- This loss function makes intuitive sense
 - If $y_n = 1$ but μ_n is close to 0 (model makes error) then loss will be high

Logistic Regression: The Loss Function

- What loss function to use? One option is to use the squared loss

$$\ell(y_n, f(\mathbf{x}_n)) = (y_n - f(\mathbf{x}_n))^2 = (y_n - \mu_n)^2 = (y_n - \sigma(\mathbf{w}^\top \mathbf{x}_n))^2$$

- This is **non-convex** and not easy to optimize
- Consider the following loss function

$$\ell(y_n, f(\mathbf{x}_n)) = \begin{cases} -\log(\mu_n) & y_n = 1 \\ -\log(1 - \mu_n) & y_n = 0 \end{cases}$$

- This loss function makes intuitive sense
 - If $y_n = 1$ but μ_n is close to 0 (model makes error) then loss will be high
 - If $y_n = 0$ but μ_n is close to 1 (model makes error) then loss will be high

Logistic Regression: The Loss Function

- What loss function to use? One option is to use the squared loss

$$\ell(y_n, f(\mathbf{x}_n)) = (y_n - f(\mathbf{x}_n))^2 = (y_n - \mu_n)^2 = (y_n - \sigma(\mathbf{w}^\top \mathbf{x}_n))^2$$

- This is **non-convex** and not easy to optimize
- Consider the following loss function

$$\ell(y_n, f(\mathbf{x}_n)) = \begin{cases} -\log(\mu_n) & y_n = 1 \\ -\log(1 - \mu_n) & y_n = 0 \end{cases}$$

- This loss function makes intuitive sense
 - If $y_n = 1$ but μ_n is close to 0 (model makes error) then loss will be high
 - If $y_n = 0$ but μ_n is close to 1 (model makes error) then loss will be high
- The above loss function can be combined and written more compactly as

$$\ell(y_n, f(\mathbf{x}_n)) = -y_n \log(\mu_n) - (1 - y_n) \log(1 - \mu_n)$$

Logistic Regression: The Loss Function

- What loss function to use? One option is to use the squared loss

$$\ell(y_n, f(\mathbf{x}_n)) = (y_n - f(\mathbf{x}_n))^2 = (y_n - \mu_n)^2 = (y_n - \sigma(\mathbf{w}^\top \mathbf{x}_n))^2$$

- This is **non-convex** and not easy to optimize
- Consider the following loss function

$$\ell(y_n, f(\mathbf{x}_n)) = \begin{cases} -\log(\mu_n) & y_n = 1 \\ -\log(1 - \mu_n) & y_n = 0 \end{cases}$$

- This loss function makes intuitive sense
 - If $y_n = 1$ but μ_n is close to 0 (model makes error) then loss will be high
 - If $y_n = 0$ but μ_n is close to 1 (model makes error) then loss will be high
- The above loss function can be combined and written more compactly as

$$\ell(y_n, f(\mathbf{x}_n)) = -y_n \log(\mu_n) - (1 - y_n) \log(1 - \mu_n)$$

- This is a function of the unknown parameter \mathbf{w} since $\mu_n = \sigma(\mathbf{w}^\top \mathbf{x}_n)$

Logistic Regression: The Loss Function

- The loss function over the entire training data

$$L(\mathbf{w}) = \sum_{n=1}^N \ell(y_n, f(\mathbf{x}_n)) = \sum_{n=1}^N [-y_n \log(\mu_n) - (1 - y_n) \log(1 - \mu_n)]$$

Logistic Regression: The Loss Function

- The loss function over the entire training data

$$L(\mathbf{w}) = \sum_{n=1}^N \ell(y_n, f(\mathbf{x}_n)) = \sum_{n=1}^N [-y_n \log(\mu_n) - (1 - y_n) \log(1 - \mu_n)]$$

- This is also known as the **cross-entropy loss**
 - Sum of the cross-entropies b/w true label y_n and predicted label prob. μ_n

Logistic Regression: The Loss Function

- The loss function over the entire training data

$$L(\mathbf{w}) = \sum_{n=1}^N \ell(y_n, f(\mathbf{x}_n)) = \sum_{n=1}^N [-y_n \log(\mu_n) - (1 - y_n) \log(1 - \mu_n)]$$

- This is also known as the **cross-entropy loss**
 - Sum of the cross-entropies b/w true label y_n and predicted label prob. μ_n
- Plugging in $\mu_n = \frac{\exp(\mathbf{w}^\top \mathbf{x}_n)}{1 + \exp(\mathbf{w}^\top \mathbf{x}_n)}$ and chugging, we get (verify yourself)

$$L(\mathbf{w}) = - \sum_{n=1}^N (y_n \mathbf{w}^\top \mathbf{x}_n - \log(1 + \exp(\mathbf{w}^\top \mathbf{x}_n)))$$

Logistic Regression: The Loss Function

- The loss function over the entire training data

$$L(\mathbf{w}) = \sum_{n=1}^N \ell(y_n, f(\mathbf{x}_n)) = \sum_{n=1}^N [-y_n \log(\mu_n) - (1 - y_n) \log(1 - \mu_n)]$$

- This is also known as the **cross-entropy loss**
 - Sum of the cross-entropies b/w true label y_n and predicted label prob. μ_n
- Plugging in $\mu_n = \frac{\exp(\mathbf{w}^\top \mathbf{x}_n)}{1 + \exp(\mathbf{w}^\top \mathbf{x}_n)}$ and chugging, we get (verify yourself)

$$L(\mathbf{w}) = - \sum_{n=1}^N (y_n \mathbf{w}^\top \mathbf{x}_n - \log(1 + \exp(\mathbf{w}^\top \mathbf{x}_n)))$$

- We can add a regularizer (e.g., squared ℓ_2 norm of \mathbf{w}) to prevent overfitting

$$L(\mathbf{w}) = - \sum_{n=1}^N (y_n \mathbf{w}^\top \mathbf{x}_n - \log(1 + \exp(\mathbf{w}^\top \mathbf{x}_n))) + \lambda \|\mathbf{w}\|^2$$

Probabilistic Modeling View (MLE/MAP) for Logistic Regression

Logistic Regression: MLE Formulation

- Recall, each label y_n is binary with prob. μ_n . Assume Bernoulli likelihood:

$$p(\mathbf{y}|\mathbf{X}, \mathbf{w}) = \prod_{n=1}^N p(y_n|\mathbf{x}_n, \mathbf{w}) = \prod_{n=1}^N \mu_n^{y_n} (1 - \mu_n)^{1-y_n}$$

where $\mu_n = \frac{\exp(\mathbf{w}^\top \mathbf{x}_n)}{1 + \exp(\mathbf{w}^\top \mathbf{x}_n)}$

Logistic Regression: MLE Formulation

- Recall, each label y_n is binary with prob. μ_n . Assume Bernoulli likelihood:

$$p(\mathbf{y}|\mathbf{X}, \mathbf{w}) = \prod_{n=1}^N p(y_n|\mathbf{x}_n, \mathbf{w}) = \prod_{n=1}^N \mu_n^{y_n} (1 - \mu_n)^{1-y_n}$$

where $\mu_n = \frac{\exp(\mathbf{w}^\top \mathbf{x}_n)}{1 + \exp(\mathbf{w}^\top \mathbf{x}_n)}$

- Doing MLE would require maximizing the log likelihood w.r.t. \mathbf{w}

$$\log p(\mathbf{Y}|\mathbf{X}, \mathbf{w}) = \sum_{n=1}^N (y_n \log \mu_n + (1 - y_n) \log(1 - \mu_n))$$

Logistic Regression: MLE Formulation

- Recall, each label y_n is binary with prob. μ_n . Assume Bernoulli likelihood:

$$p(\mathbf{y}|\mathbf{X}, \mathbf{w}) = \prod_{n=1}^N p(y_n|\mathbf{x}_n, \mathbf{w}) = \prod_{n=1}^N \mu_n^{y_n} (1 - \mu_n)^{1-y_n}$$

$$\text{where } \mu_n = \frac{\exp(\mathbf{w}^\top \mathbf{x}_n)}{1 + \exp(\mathbf{w}^\top \mathbf{x}_n)}$$

- Doing MLE would require maximizing the log likelihood w.r.t. \mathbf{w}

$$\log p(\mathbf{Y}|\mathbf{X}, \mathbf{w}) = \sum_{n=1}^N (y_n \log \mu_n + (1 - y_n) \log(1 - \mu_n))$$

- This is equivalent to minimizing the NLL. Plugging in $\mu_n = \frac{\exp(\mathbf{w}^\top \mathbf{x}_n)}{1 + \exp(\mathbf{w}^\top \mathbf{x}_n)}$ we get

$$\text{NLL}(\mathbf{w}) = - \sum_{n=1}^N (y_n \mathbf{w}^\top \mathbf{x}_n - \log(1 + \exp(\mathbf{w}^\top \mathbf{x}_n)))$$

Logistic Regression: MLE Formulation

- Recall, each label y_n is binary with prob. μ_n . Assume Bernoulli likelihood:

$$p(\mathbf{y}|\mathbf{X}, \mathbf{w}) = \prod_{n=1}^N p(y_n|\mathbf{x}_n, \mathbf{w}) = \prod_{n=1}^N \mu_n^{y_n} (1 - \mu_n)^{1-y_n}$$

$$\text{where } \mu_n = \frac{\exp(\mathbf{w}^\top \mathbf{x}_n)}{1 + \exp(\mathbf{w}^\top \mathbf{x}_n)}$$

- Doing MLE would require maximizing the log likelihood w.r.t. \mathbf{w}

$$\log p(\mathbf{Y}|\mathbf{X}, \mathbf{w}) = \sum_{n=1}^N (y_n \log \mu_n + (1 - y_n) \log(1 - \mu_n))$$

- This is equivalent to minimizing the NLL. Plugging in $\mu_n = \frac{\exp(\mathbf{w}^\top \mathbf{x}_n)}{1 + \exp(\mathbf{w}^\top \mathbf{x}_n)}$ we get

$$\text{NLL}(\mathbf{w}) = - \sum_{n=1}^N (y_n \mathbf{w}^\top \mathbf{x}_n - \log(1 + \exp(\mathbf{w}^\top \mathbf{x}_n)))$$

- Not surprisingly, the NLL expression is the same as the loss function

Logistic Regression: MAP Formulation

- MLE estimate of \mathbf{w} can lead to overfitting. Solution: use a prior on \mathbf{w}
- Just like the linear regression case, let's put a Gaussian prior on \mathbf{w}

$$p(\mathbf{w}) = \mathcal{N}(0, \lambda^{-1} \mathbf{I}_D) \propto \exp(-\lambda \mathbf{w}^\top \mathbf{w})$$

Logistic Regression: MAP Formulation

- MLE estimate of \mathbf{w} can lead to overfitting. Solution: use a prior on \mathbf{w}
- Just like the linear regression case, let's put a Gaussian prior on \mathbf{w}

$$p(\mathbf{w}) = \mathcal{N}(0, \lambda^{-1} \mathbf{I}_D) \propto \exp(-\lambda \mathbf{w}^\top \mathbf{w})$$

- MAP objective: MLE objective + $\log p(\mathbf{w})$

Logistic Regression: MAP Formulation

- MLE estimate of \mathbf{w} can lead to overfitting. Solution: use a prior on \mathbf{w}
- Just like the linear regression case, let's put a Gaussian prior on \mathbf{w}

$$p(\mathbf{w}) = \mathcal{N}(0, \lambda^{-1} \mathbf{I}_D) \propto \exp(-\lambda \mathbf{w}^\top \mathbf{w})$$

- MAP objective: MLE objective + $\log p(\mathbf{w})$
- Can maximize the MAP objective (log-posterior) w.r.t. \mathbf{w} or minimize the negative of log posterior which will be

$$\text{NLL}(\mathbf{w}) - \log p(\mathbf{w})$$

Logistic Regression: MAP Formulation

- MLE estimate of \mathbf{w} can lead to overfitting. Solution: use a prior on \mathbf{w}
- Just like the linear regression case, let's put a Gaussian prior on \mathbf{w}

$$p(\mathbf{w}) = \mathcal{N}(0, \lambda^{-1} \mathbf{I}_D) \propto \exp(-\lambda \mathbf{w}^\top \mathbf{w})$$

- MAP objective: MLE objective + $\log p(\mathbf{w})$
- Can maximize the MAP objective (log-posterior) w.r.t. \mathbf{w} or minimize the negative of log posterior which will be

$$\text{NLL}(\mathbf{w}) - \log p(\mathbf{w})$$

- Ignoring the constants, we get the following objective for MAP estimation

$$-\sum_{n=1}^N (y_n \mathbf{w}^\top \mathbf{x}_n - \log(1 + \exp(\mathbf{w}^\top \mathbf{x}_n))) + \lambda \mathbf{w}^\top \mathbf{w}$$

Logistic Regression: MAP Formulation

- MLE estimate of \mathbf{w} can lead to overfitting. Solution: use a prior on \mathbf{w}

- Just like the linear regression case, let's put a Gaussian prior on \mathbf{w}

$$p(\mathbf{w}) = \mathcal{N}(0, \lambda^{-1} \mathbf{I}_D) \propto \exp(-\lambda \mathbf{w}^\top \mathbf{w})$$

- MAP objective: MLE objective + $\log p(\mathbf{w})$

- Can maximize the MAP objective (log-posterior) w.r.t. \mathbf{w} or minimize the negative of log posterior which will be

$$\text{NLL}(\mathbf{w}) - \log p(\mathbf{w})$$

- Ignoring the constants, we get the following objective for MAP estimation

$$-\sum_{n=1}^N (y_n \mathbf{w}^\top \mathbf{x}_n - \log(1 + \exp(\mathbf{w}^\top \mathbf{x}_n))) + \lambda \mathbf{w}^\top \mathbf{w}$$

- Thus MAP estimation is equivalent to **regularized logistic regression**

Estimating the Weight Vector \mathbf{w}

- Loss function/NLL for logistic regression (ignoring the regularizer term)

$$L(\mathbf{w}) = - \sum_{n=1}^N (y_n \mathbf{w}^\top \mathbf{x}_n - \log(1 + \exp(\mathbf{w}^\top \mathbf{x}_n)))$$

- The loss function is convex in \mathbf{w} (thus has a unique minimum)

Estimating the Weight Vector \mathbf{w}

- Loss function/NLL for logistic regression (ignoring the regularizer term)

$$L(\mathbf{w}) = - \sum_{n=1}^N (y_n \mathbf{w}^\top \mathbf{x}_n - \log(1 + \exp(\mathbf{w}^\top \mathbf{x}_n)))$$

- The loss function is convex in \mathbf{w} (thus has a unique minimum)
- The gradient/derivative of $L(\mathbf{w})$ w.r.t. \mathbf{w} (let's ignore the regularizer)

$$\mathbf{g} = \frac{\partial L(\mathbf{w})}{\partial \mathbf{w}} = \frac{\partial}{\partial \mathbf{w}} \left[- \sum_{n=1}^N (y_n \mathbf{w}^\top \mathbf{x}_n - \log(1 + \exp(\mathbf{w}^\top \mathbf{x}_n))) \right]$$

Estimating the Weight Vector \mathbf{w}

- Loss function/NLL for logistic regression (ignoring the regularizer term)

$$L(\mathbf{w}) = - \sum_{n=1}^N (y_n \mathbf{w}^\top \mathbf{x}_n - \log(1 + \exp(\mathbf{w}^\top \mathbf{x}_n)))$$

- The loss function is convex in \mathbf{w} (thus has a unique minimum)
- The gradient/derivative of $L(\mathbf{w})$ w.r.t. \mathbf{w} (let's ignore the regularizer)

$$\begin{aligned} \mathbf{g} = \frac{\partial L(\mathbf{w})}{\partial \mathbf{w}} &= \frac{\partial}{\partial \mathbf{w}} \left[- \sum_{n=1}^N (y_n \mathbf{w}^\top \mathbf{x}_n - \log(1 + \exp(\mathbf{w}^\top \mathbf{x}_n))) \right] \\ &= - \sum_{n=1}^N \left(y_n \mathbf{x}_n - \frac{\exp(\mathbf{w}^\top \mathbf{x}_n)}{(1 + \exp(\mathbf{w}^\top \mathbf{x}_n))} \mathbf{x}_n \right) \end{aligned}$$

Estimating the Weight Vector \mathbf{w}

- Loss function/NLL for logistic regression (ignoring the regularizer term)

$$L(\mathbf{w}) = - \sum_{n=1}^N (y_n \mathbf{w}^\top \mathbf{x}_n - \log(1 + \exp(\mathbf{w}^\top \mathbf{x}_n)))$$

- The loss function is convex in \mathbf{w} (thus has a unique minimum)
- The gradient/derivative of $L(\mathbf{w})$ w.r.t. \mathbf{w} (let's ignore the regularizer)

$$\begin{aligned} \mathbf{g} = \frac{\partial L(\mathbf{w})}{\partial \mathbf{w}} &= \frac{\partial}{\partial \mathbf{w}} \left[- \sum_{n=1}^N (y_n \mathbf{w}^\top \mathbf{x}_n - \log(1 + \exp(\mathbf{w}^\top \mathbf{x}_n))) \right] \\ &= - \sum_{n=1}^N \left(y_n \mathbf{x}_n - \frac{\exp(\mathbf{w}^\top \mathbf{x}_n)}{(1 + \exp(\mathbf{w}^\top \mathbf{x}_n))} \mathbf{x}_n \right) \\ &= - \sum_{n=1}^N (y_n - \mu_n) \mathbf{x}_n = \mathbf{X}^\top (\boldsymbol{\mu} - \mathbf{y}) \end{aligned}$$

Estimating the Weight Vector \mathbf{w}

- Loss function/NLL for logistic regression (ignoring the regularizer term)

$$L(\mathbf{w}) = - \sum_{n=1}^N (y_n \mathbf{w}^\top \mathbf{x}_n - \log(1 + \exp(\mathbf{w}^\top \mathbf{x}_n)))$$

- The loss function is convex in \mathbf{w} (thus has a unique minimum)
- The gradient/derivative of $L(\mathbf{w})$ w.r.t. \mathbf{w} (let's ignore the regularizer)

$$\begin{aligned} \mathbf{g} = \frac{\partial L(\mathbf{w})}{\partial \mathbf{w}} &= \frac{\partial}{\partial \mathbf{w}} \left[- \sum_{n=1}^N (y_n \mathbf{w}^\top \mathbf{x}_n - \log(1 + \exp(\mathbf{w}^\top \mathbf{x}_n))) \right] \\ &= - \sum_{n=1}^N \left(y_n \mathbf{x}_n - \frac{\exp(\mathbf{w}^\top \mathbf{x}_n)}{(1 + \exp(\mathbf{w}^\top \mathbf{x}_n))} \mathbf{x}_n \right) \\ &= - \sum_{n=1}^N (y_n - \mu_n) \mathbf{x}_n = \mathbf{X}^\top (\boldsymbol{\mu} - \mathbf{y}) \end{aligned}$$

- Can't get a closed form solution for \mathbf{w} by setting the derivative to zero
 - Need to use iterative methods (e.g., gradient descent) to solve for \mathbf{w}

Gradient Descent for Logistic Regression

- We can use gradient descent (GD) to solve for \mathbf{w} as follows:

Gradient Descent for Logistic Regression

- We can use gradient descent (GD) to solve for \mathbf{w} as follows:
 - Initialize $\mathbf{w}^{(1)} \in \mathbb{R}^D$ randomly.

Gradient Descent for Logistic Regression

- We can use gradient descent (GD) to solve for \mathbf{w} as follows:
 - Initialize $\mathbf{w}^{(1)} \in \mathbb{R}^D$ randomly.
 - Iterate the following until convergence

$$\underbrace{\mathbf{w}^{(t+1)}}_{\text{new value}} = \underbrace{\mathbf{w}^{(t)}}_{\text{previous value}} - \eta \underbrace{\sum_{n=1}^N (\mu_n^{(t)} - y_n) \mathbf{x}_n}_{\text{gradient at previous value}}$$

where η is the **learning rate** and $\mu^{(t)} = \sigma(\mathbf{w}^{(t)\top} \mathbf{x}_n)$ is the predicted label probability for \mathbf{x}_n using $\mathbf{w} = \mathbf{w}^{(t)}$ from the previous iteration

Gradient Descent for Logistic Regression

- We can use gradient descent (GD) to solve for \mathbf{w} as follows:
 - Initialize $\mathbf{w}^{(1)} \in \mathbb{R}^D$ randomly.
 - Iterate the following until convergence

$$\underbrace{\mathbf{w}^{(t+1)}}_{\text{new value}} = \underbrace{\mathbf{w}^{(t)}}_{\text{previous value}} - \eta \underbrace{\sum_{n=1}^N (\mu_n^{(t)} - y_n) \mathbf{x}_n}_{\text{gradient at previous value}}$$

where η is the **learning rate** and $\mu^{(t)} = \sigma(\mathbf{w}^{(t)\top} \mathbf{x}_n)$ is the predicted label probability for \mathbf{x}_n using $\mathbf{w} = \mathbf{w}^{(t)}$ from the previous iteration

- Note that the updates give larger weights to those examples on which the current model makes larger mistakes, as measured by $(\mu_n^{(t)} - y_n)$

Gradient Descent for Logistic Regression

- We can use gradient descent (GD) to solve for \mathbf{w} as follows:
 - Initialize $\mathbf{w}^{(1)} \in \mathbb{R}^D$ randomly.
 - Iterate the following until convergence

$$\underbrace{\mathbf{w}^{(t+1)}}_{\text{new value}} = \underbrace{\mathbf{w}^{(t)}}_{\text{previous value}} - \eta \underbrace{\sum_{n=1}^N (\mu_n^{(t)} - y_n) \mathbf{x}_n}_{\text{gradient at previous value}}$$

where η is the **learning rate** and $\mu^{(t)} = \sigma(\mathbf{w}^{(t)\top} \mathbf{x}_n)$ is the predicted label probability for \mathbf{x}_n using $\mathbf{w} = \mathbf{w}^{(t)}$ from the previous iteration

- Note that the updates give larger weights to those examples on which the current model makes larger mistakes, as measured by $(\mu_n^{(t)} - y_n)$
- **Note:** Computing the gradient in every iteration requires all the data. Thus GD can be expensive if N is very large. A cheaper alternative is to do GD using only a **small randomly chosen minibatch** of data. It is known as **Stochastic Gradient Descent** (SGD). Runs faster and converges faster.

More on Gradient Descent..

- GD can converge slowly and is also sensitive to the step size




Figure: Left: small step sizes. Right: large step sizes

¹Also see: "A comparison of numerical optimizers for logistic regression" by Tom Minka

More on Gradient Descent..

- GD can converge slowly and is also sensitive to the step size




Figure: Left: small step sizes. Right: large step sizes

- Several ways to remedy this¹.

¹Also see: "A comparison of numerical optimizers for logistic regression" by Tom Minka

More on Gradient Descent..

- GD can converge slowly and is also sensitive to the step size




Figure: Left: small step sizes. Right: large step sizes

- Several ways to remedy this¹. E.g.,

¹Also see: "A comparison of numerical optimizers for logistic regression" by Tom Minka

More on Gradient Descent..

- GD can converge slowly and is also sensitive to the step size




Figure: Left: small step sizes. Right: large step sizes

- Several ways to remedy this¹. E.g.,
 - Choose the optimal step size η_t (different in each iteration) by **line-search**

¹Also see: "A comparison of numerical optimizers for logistic regression" by Tom Minka

More on Gradient Descent..

- GD can converge slowly and is also sensitive to the step size




Figure: Left: small step sizes. Right: large step sizes

- Several ways to remedy this¹. E.g.,
 - Choose the optimal step size η_t (different in each iteration) by **line-search**
 - Add a **momentum term** to the updates

$$\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} - \eta_t \mathbf{g}^{(t)} + \alpha_t (\mathbf{w}^{(t)} - \mathbf{w}^{(t-1)})$$

¹Also see: "A comparison of numerical optimizers for logistic regression" by Tom Minka

Multiclass Logistic (or “Softmax”) Regression

- Logistic regression can be extended to handle $K > 2$ classes
- In this case, $y_n \in \{0, 1, 2, \dots, K - 1\}$ and label probabilities are defined as

$$p(y_n = k | \mathbf{x}_n, \mathbf{W}) = \frac{\exp(\mathbf{w}_k^\top \mathbf{x}_n)}{\sum_{\ell=1}^K \exp(\mathbf{w}_\ell^\top \mathbf{x}_n)} = \mu_{nk}$$

- μ_{nk} : probability that example n belongs to class k . Also, $\sum_{\ell=1}^K \mu_{n\ell} = 1$
- $\mathbf{W} = [\mathbf{w}_1 \ \mathbf{w}_2 \ \dots \ \mathbf{w}_K]$ is $D \times K$ **weight matrix** (column k for class k)

Multiclass Logistic (or “Softmax”) Regression

- Logistic regression can be extended to handle $K > 2$ classes
- In this case, $y_n \in \{0, 1, 2, \dots, K - 1\}$ and label probabilities are defined as

$$p(y_n = k | \mathbf{x}_n, \mathbf{W}) = \frac{\exp(\mathbf{w}_k^\top \mathbf{x}_n)}{\sum_{\ell=1}^K \exp(\mathbf{w}_\ell^\top \mathbf{x}_n)} = \mu_{nk}$$

- μ_{nk} : probability that example n belongs to class k . Also, $\sum_{\ell=1}^K \mu_{n\ell} = 1$
- $\mathbf{W} = [\mathbf{w}_1 \ \mathbf{w}_2 \ \dots \ \mathbf{w}_K]$ is $D \times K$ **weight matrix** (column k for class k)
- “Softmax” because class ‘ k ’ with largest $\mathbf{w}_k^\top \mathbf{x}_n$ dominates the probability

Multiclass Logistic (or “Softmax”) Regression

- Logistic regression can be extended to handle $K > 2$ classes
- In this case, $y_n \in \{0, 1, 2, \dots, K - 1\}$ and label probabilities are defined as

$$p(y_n = k | \mathbf{x}_n, \mathbf{W}) = \frac{\exp(\mathbf{w}_k^\top \mathbf{x}_n)}{\sum_{\ell=1}^K \exp(\mathbf{w}_\ell^\top \mathbf{x}_n)} = \mu_{nk}$$

- μ_{nk} : probability that example n belongs to class k . Also, $\sum_{\ell=1}^K \mu_{n\ell} = 1$
- $\mathbf{W} = [\mathbf{w}_1 \ \mathbf{w}_2 \ \dots \ \mathbf{w}_K]$ is $D \times K$ **weight matrix** (column k for class k)
- “Softmax” because class ‘ k ’ with largest $\mathbf{w}_k^\top \mathbf{x}_n$ dominates the probability
- We can think of the y_n ’s as drawn from a multinomial distribution

$$p(\mathbf{y} | \mathbf{X}, \mathbf{W}) = \prod_{n=1}^N \prod_{\ell=1}^K \mu_{n\ell}^{y_{n\ell}} \quad (\text{Likelihood function})$$

where $y_{n\ell} = 1$ if true class of example n is ℓ and $y_{n\ell'} = 0$ for all other $\ell' \neq \ell$

Multiclass Logistic (or “Softmax”) Regression

- Logistic regression can be extended to handle $K > 2$ classes
- In this case, $y_n \in \{0, 1, 2, \dots, K - 1\}$ and label probabilities are defined as

$$p(y_n = k | \mathbf{x}_n, \mathbf{W}) = \frac{\exp(\mathbf{w}_k^\top \mathbf{x}_n)}{\sum_{\ell=1}^K \exp(\mathbf{w}_\ell^\top \mathbf{x}_n)} = \mu_{nk}$$

- μ_{nk} : probability that example n belongs to class k . Also, $\sum_{\ell=1}^K \mu_{n\ell} = 1$
- $\mathbf{W} = [\mathbf{w}_1 \ \mathbf{w}_2 \ \dots \ \mathbf{w}_K]$ is $D \times K$ **weight matrix** (column k for class k)
- “Softmax” because class ‘ k ’ with largest $\mathbf{w}_k^\top \mathbf{x}_n$ dominates the probability
- We can think of the y_n ’s as drawn from a multinomial distribution

$$p(\mathbf{y} | \mathbf{X}, \mathbf{W}) = \prod_{n=1}^N \prod_{\ell=1}^K \mu_{n\ell}^{y_{n\ell}} \quad (\text{Likelihood function})$$

where $y_{n\ell} = 1$ if true class of example n is ℓ and $y_{n\ell'} = 0$ for all other $\ell' \neq \ell$

- Can do MLE/MAP for \mathbf{W} similar to the binary logistic regression case

Logistic Regression: Summary

- A probabilistic model for binary classification
- Simple objective, easy to optimize using gradient based methods
- Very widely used, very efficient solvers exist
- Can be extended for multiclass (softmax) classification
- Used as modules in more complex models (e.g, deep neural nets)